Round 2 | Smaller mean gradients, persistently large noise

V4: discounted-return score with a leave-one-out baseline | 1,400 batches per policy

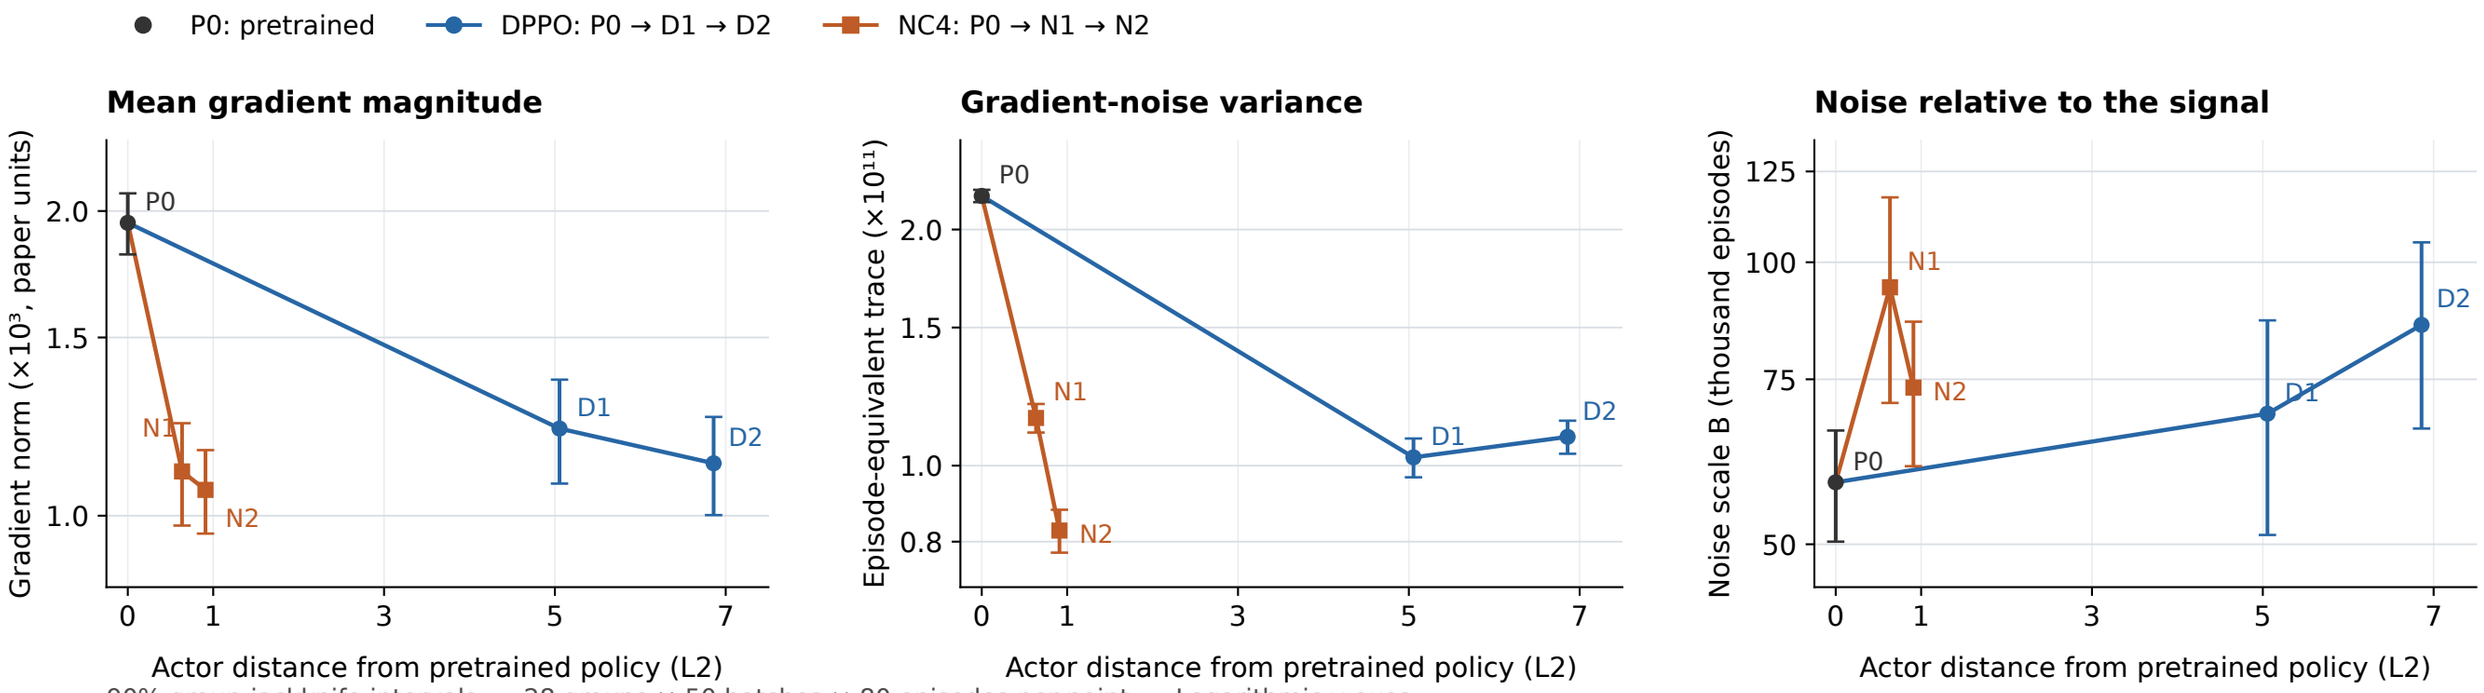

90% group-jackknife intervals • 28 groups × 50 batches × 80 episodes per point • Logarithmic y-axes

Each path uses training seed 0. Connecting lines join saved checkpoints; uncertainty is conditional on those checkpoints.