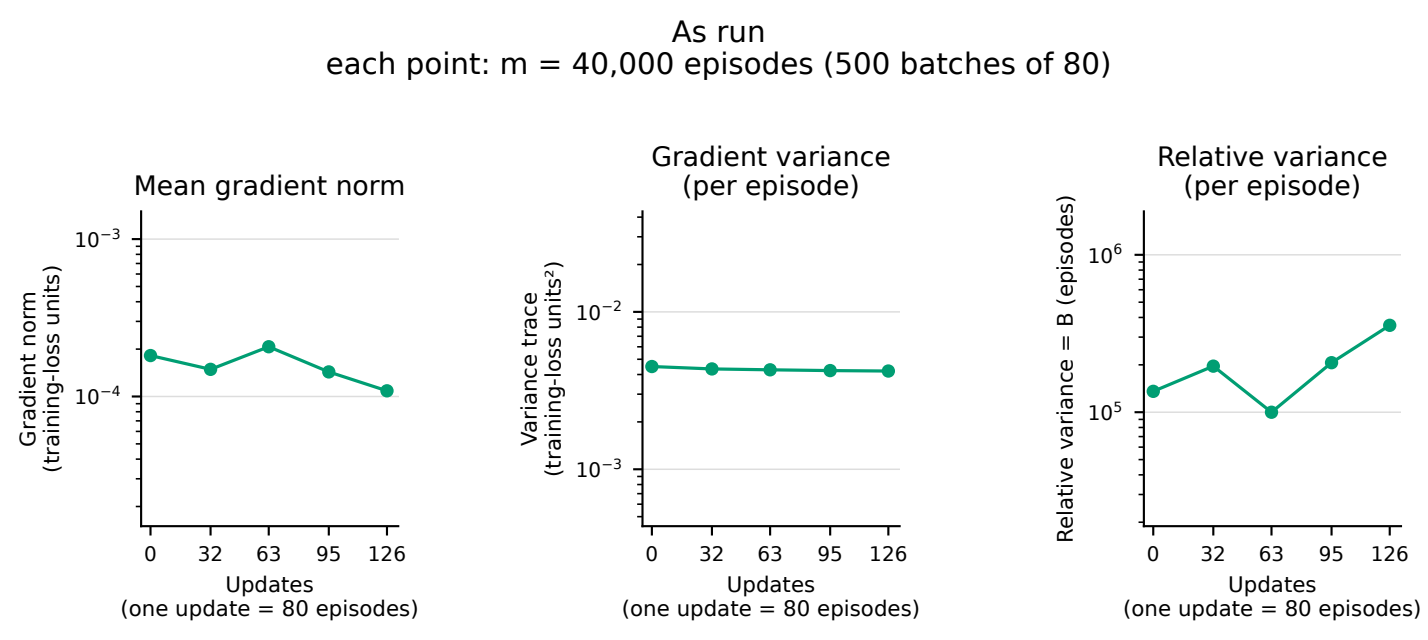
As run
each point: m = 40,000 episodes (500 batches of 80)
Mean gradient norm
Gradient norm
(training-loss units)
10^{-3}
10^{-4}
0
32
63
95
126
Updates
(one update = 80 episodes)
Gradient variance
(per episode)
Variance trace
(training-loss units²)
10^{-2}
10^{-3}
0
32
63
95
126
Updates
(one update = 80 episodes)
Relative variance
(per episode)
Relative variance = B (episodes)
10^{6}
10^{5}
0
32
63
95
126
Updates
(one update = 80 episodes)