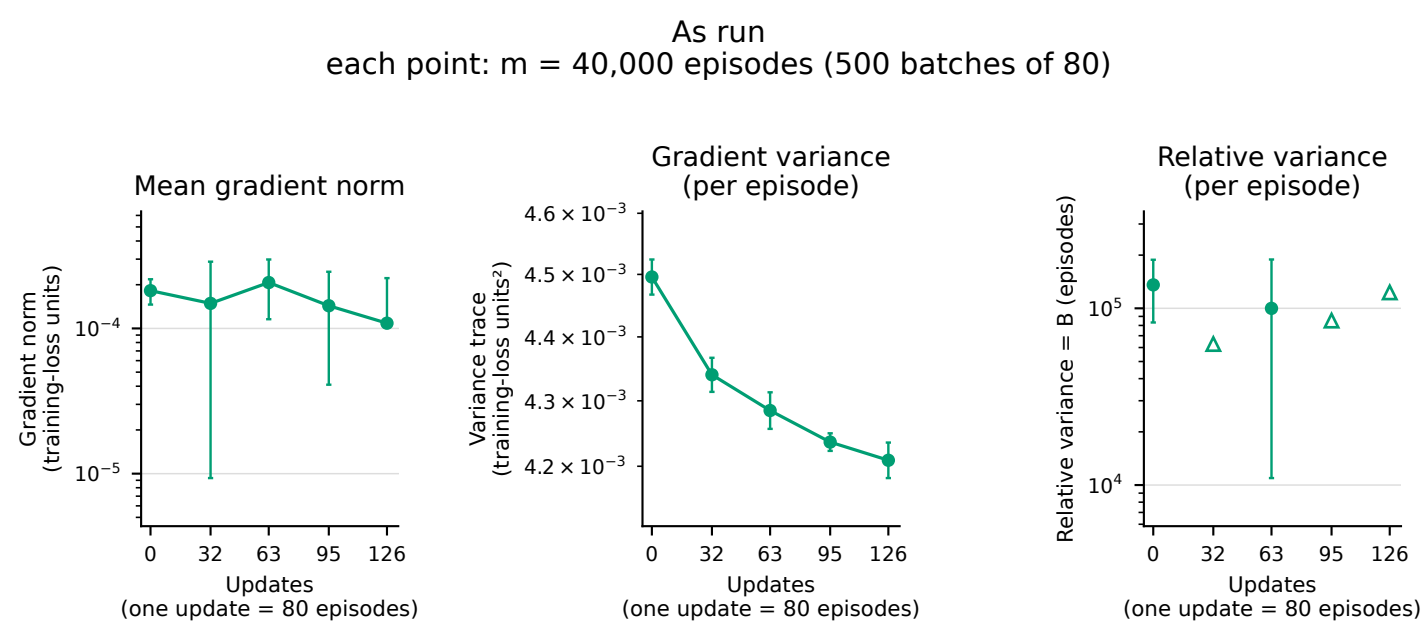
As run
each point: m = 40,000 episodes (500 batches of 80)
Mean gradient norm
Gradient norm
(training-loss units)
10^{-4}
10^{-5}
0
32
63
95
126
Updates
(one update = 80 episodes)
Gradient variance
(per episode)
Variance trace
(training-loss units²)
4.6×10^{-3}
4.5×10^{-3}
4.4×10^{-3}
4.3×10^{-3}
4.2×10^{-3}
0
32
63
95
126
Updates
(one update = 80 episodes)
Relative variance
(per episode)
Relative variance = B (episodes)
10^{5}
10^{4}
0
32
63
95
126
Updates
(one update = 80 episodes)