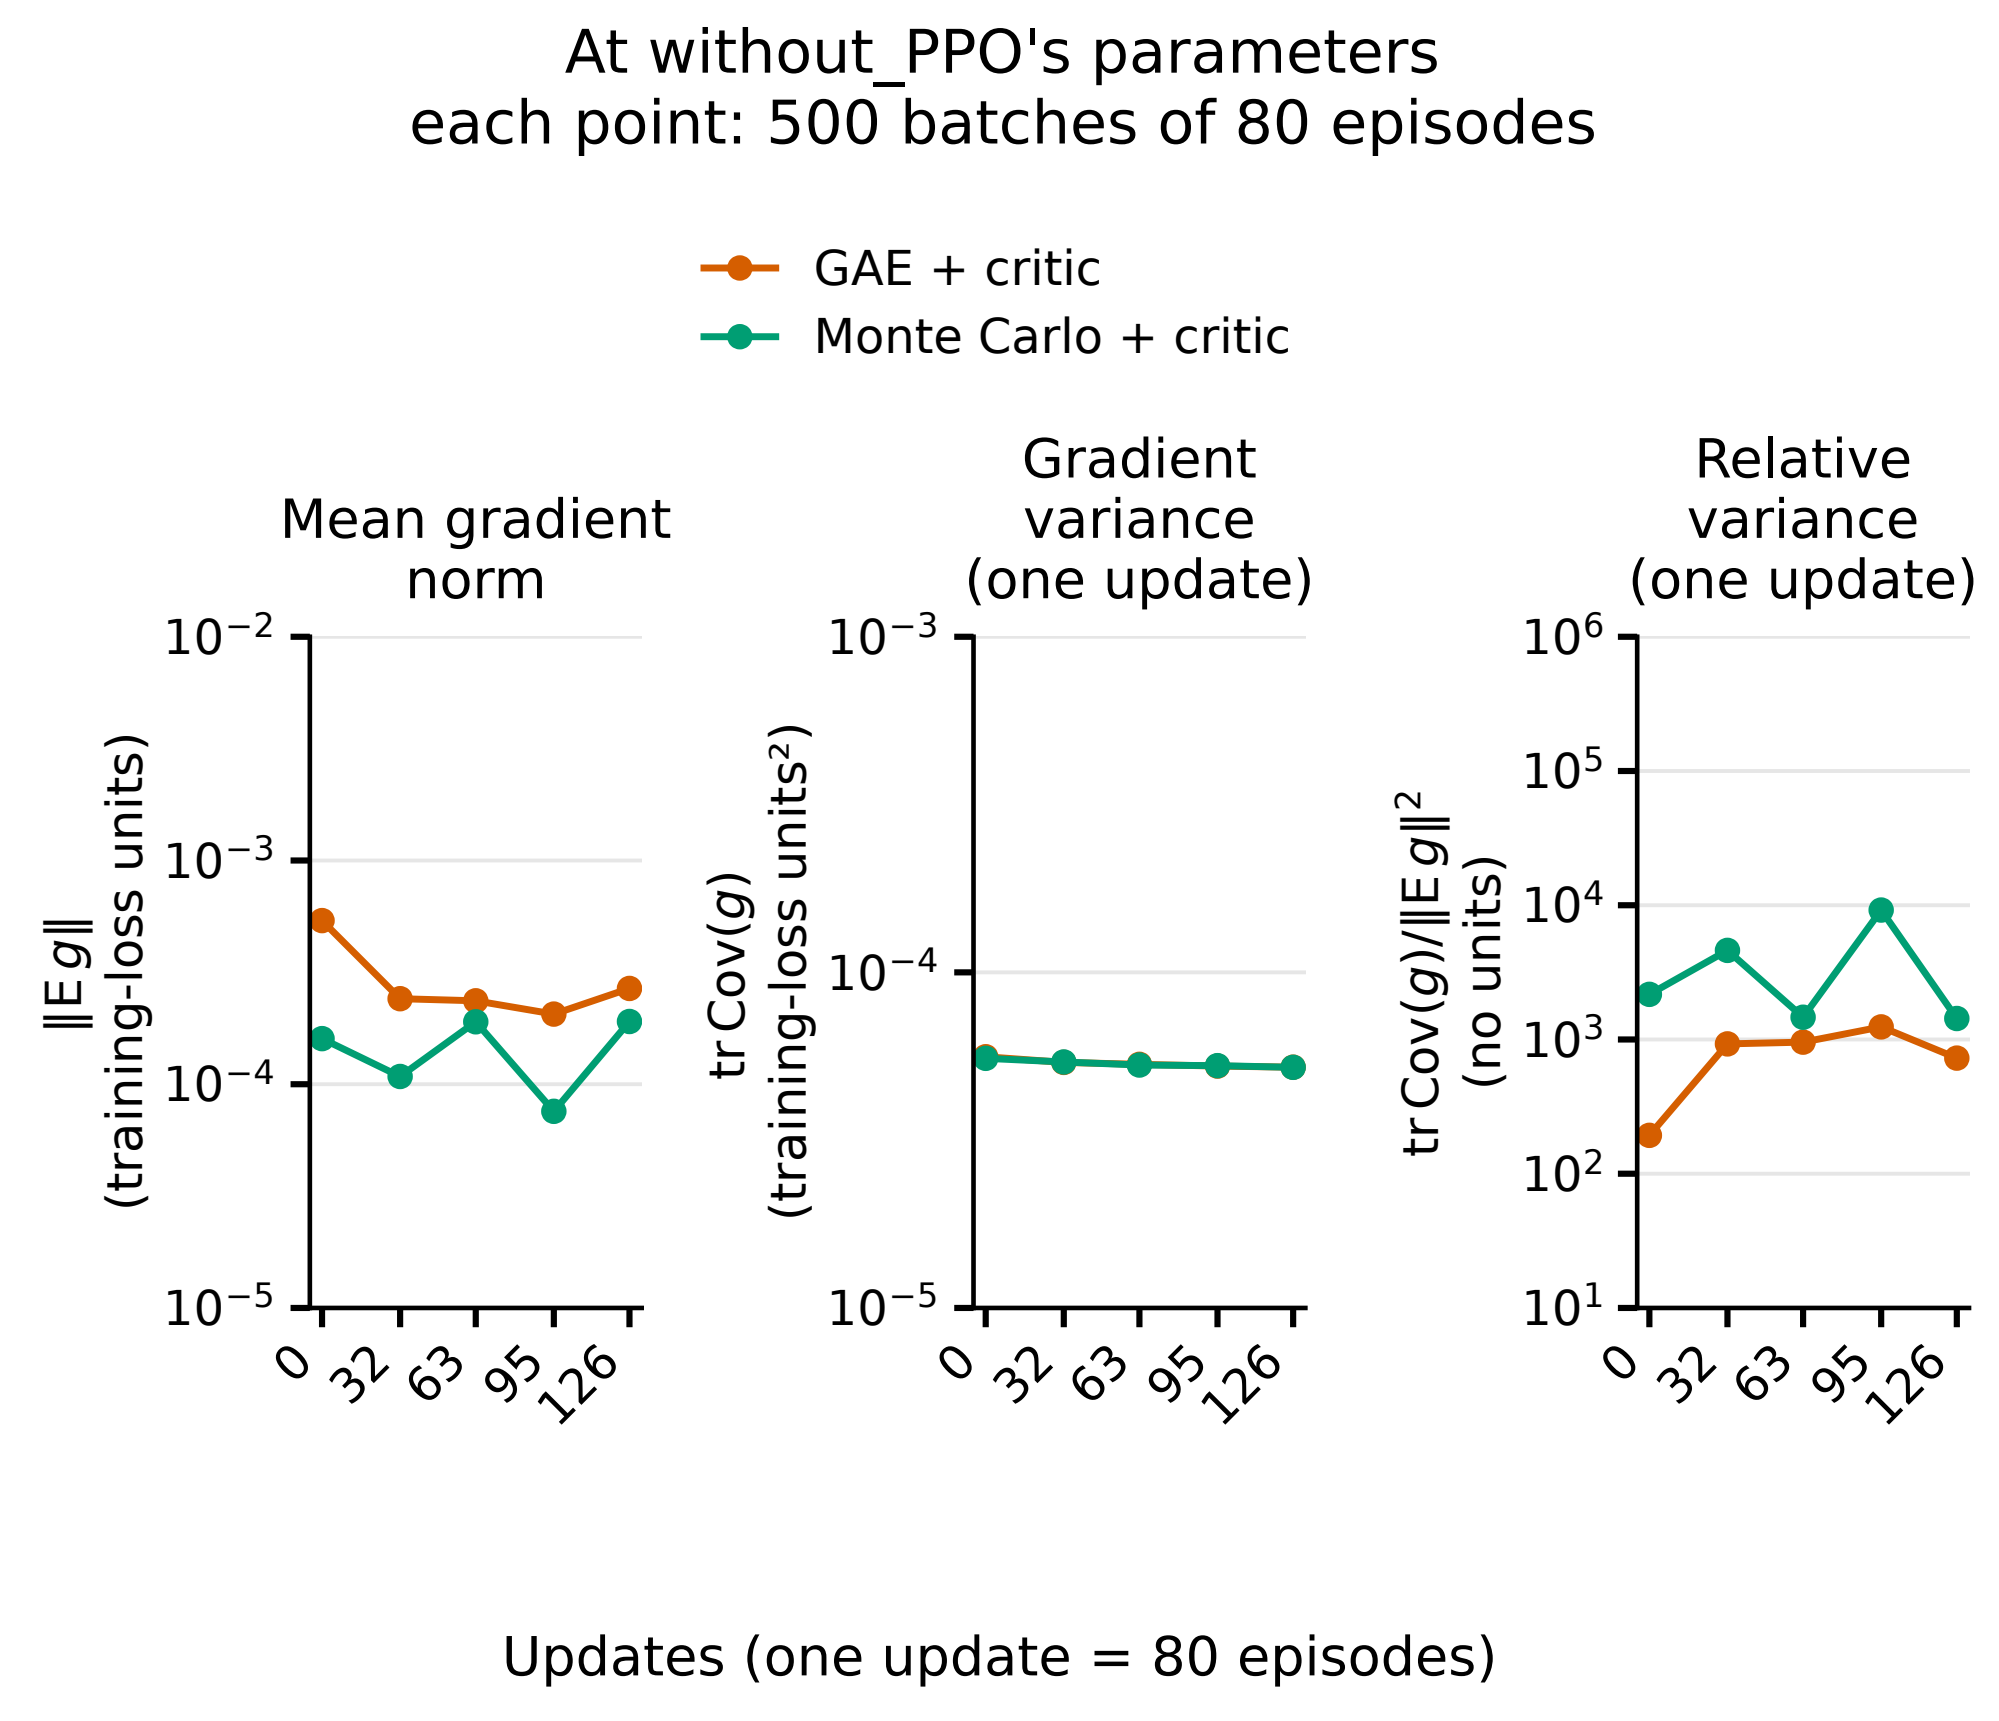
At without_PPO's parameters
each point: 500 batches of 80 episodes
GAE + critic
Monte Carlo + critic
Mean gradient norm
‖E g‖
(training-loss units)
Gradient variance (one update)
tr Cov(g)
(training-loss units²)
Relative variance (one update)
tr Cov(g)/‖E g‖²
(no units)
Updates (one update = 80 episodes)