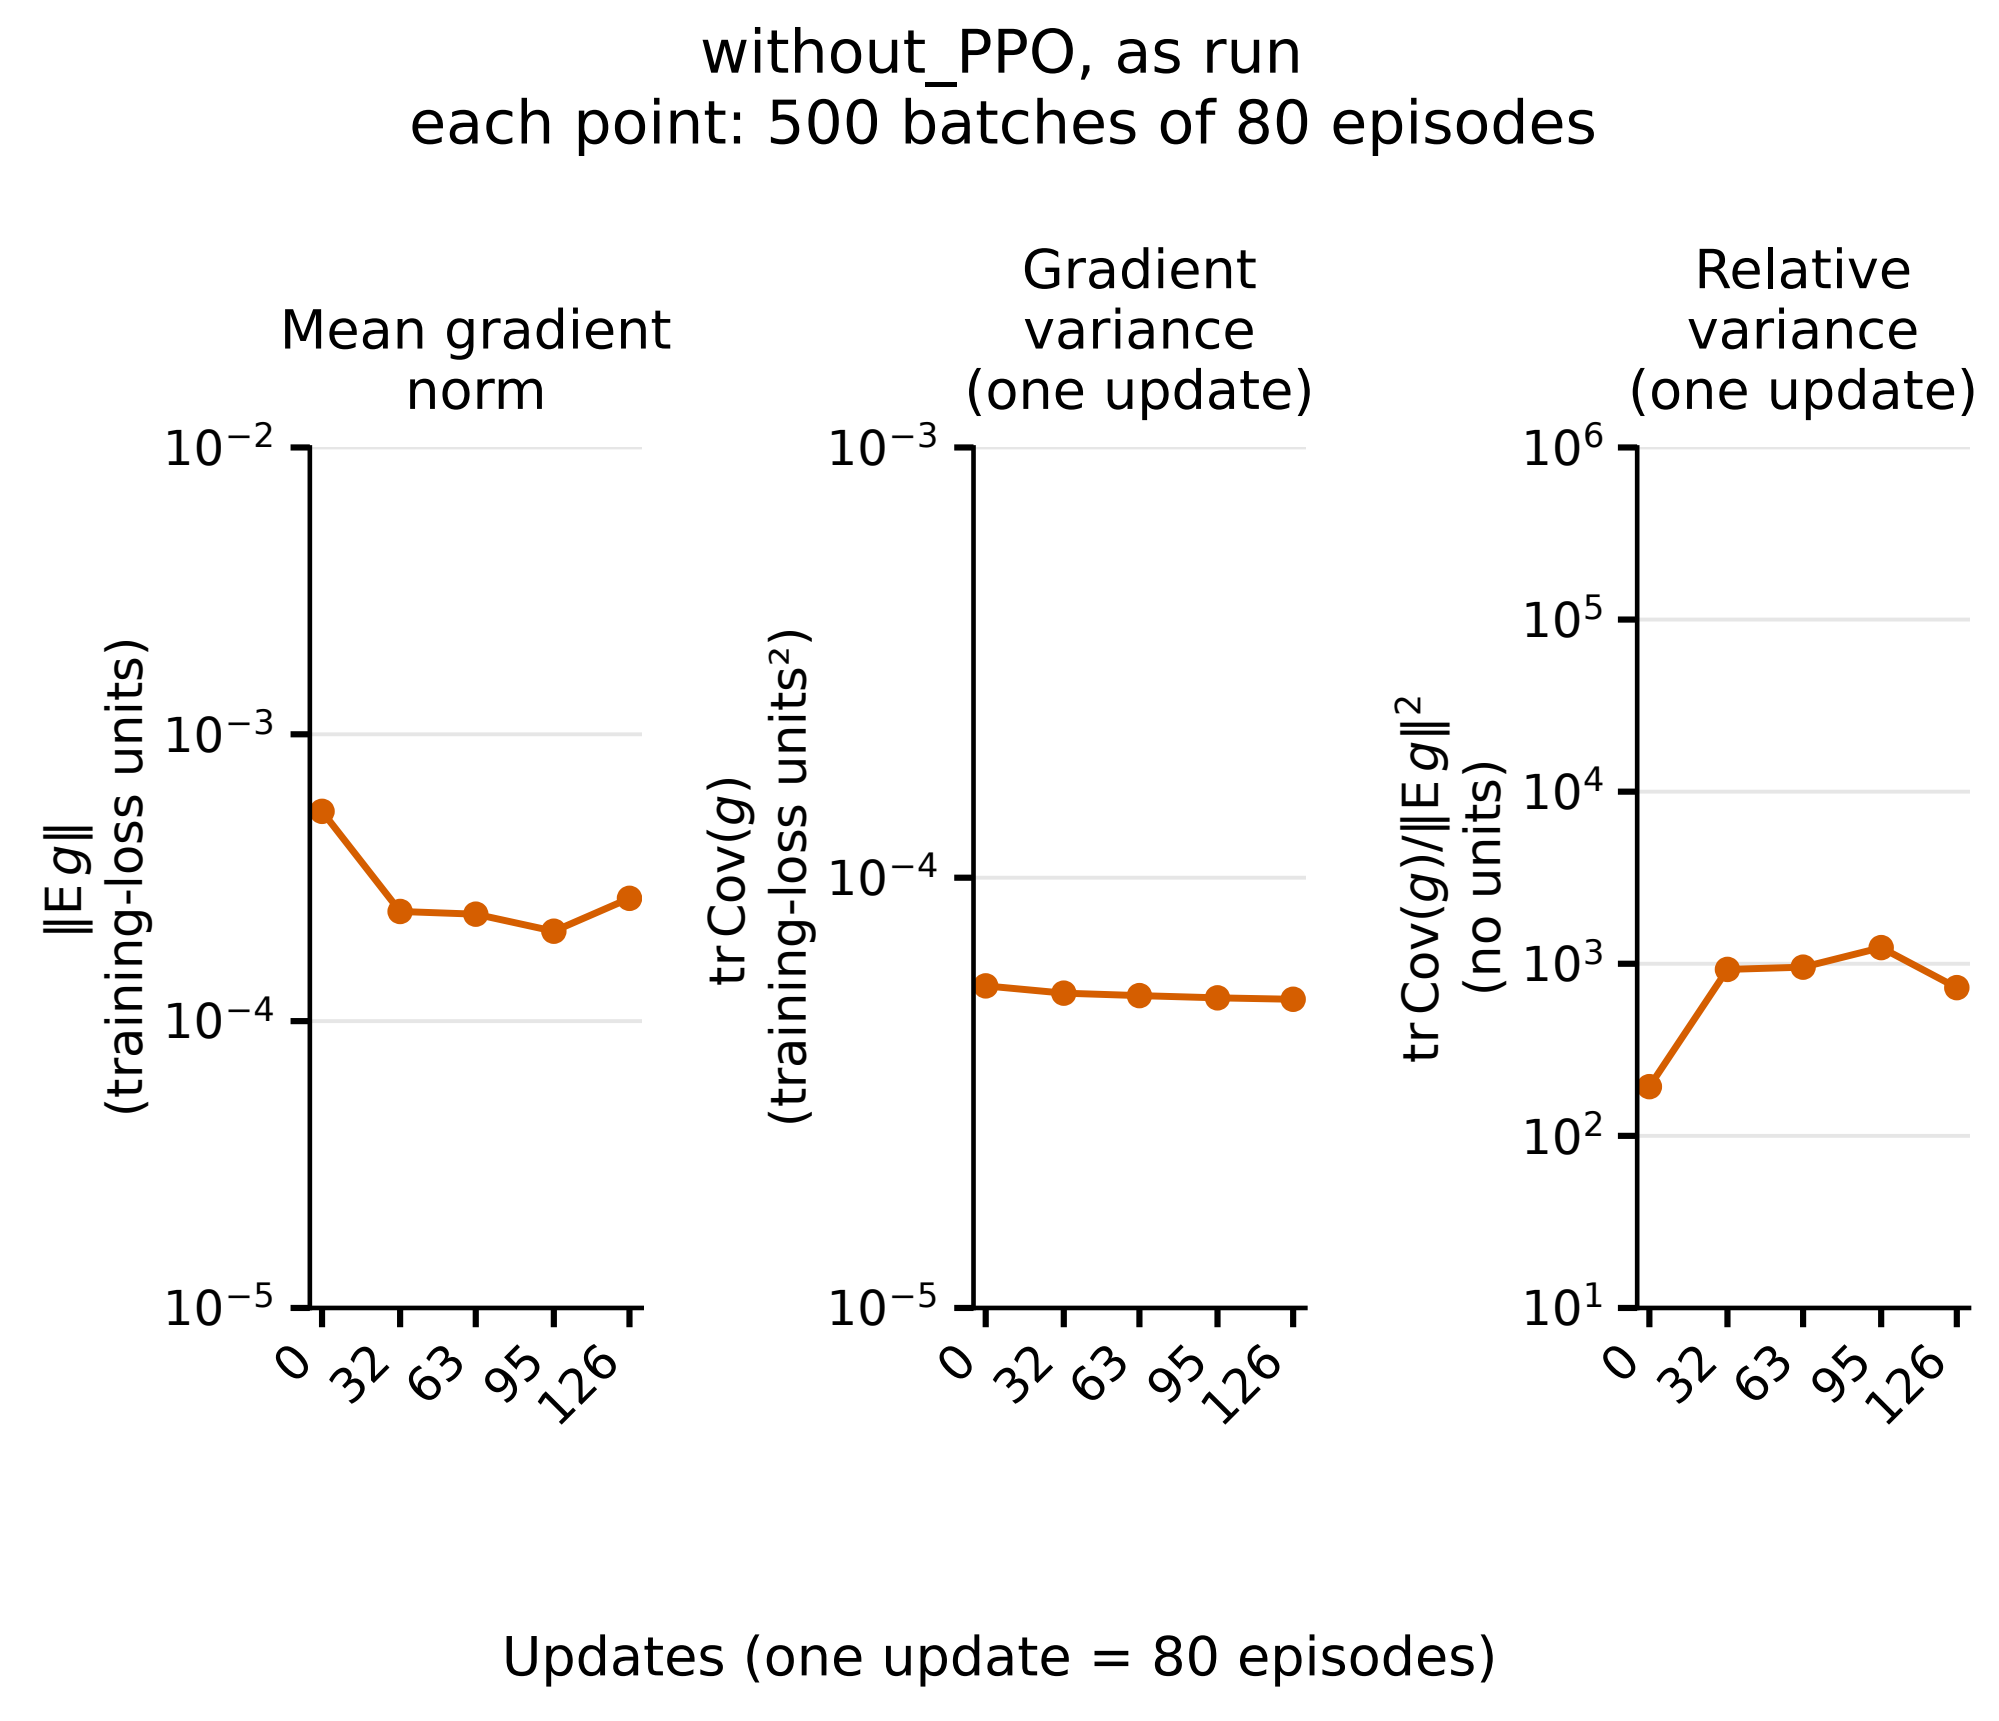
without_PPO, as run
each point: 500 batches of 80 episodes
Mean gradient norm
Gradient variance (one update)
Relative variance (one update)
‖E g‖ (training-loss units)
tr Cov(g) (training-loss units²)
tr Cov(g)/‖E g‖² (no units)
10^−2
10^−3
10^−4
10^−5
10^−3
10^−4
10^−5
10^6
10^5
10^4
10^3
10^2
10^1
0
32
63
95
126
Updates (one update = 80 episodes)