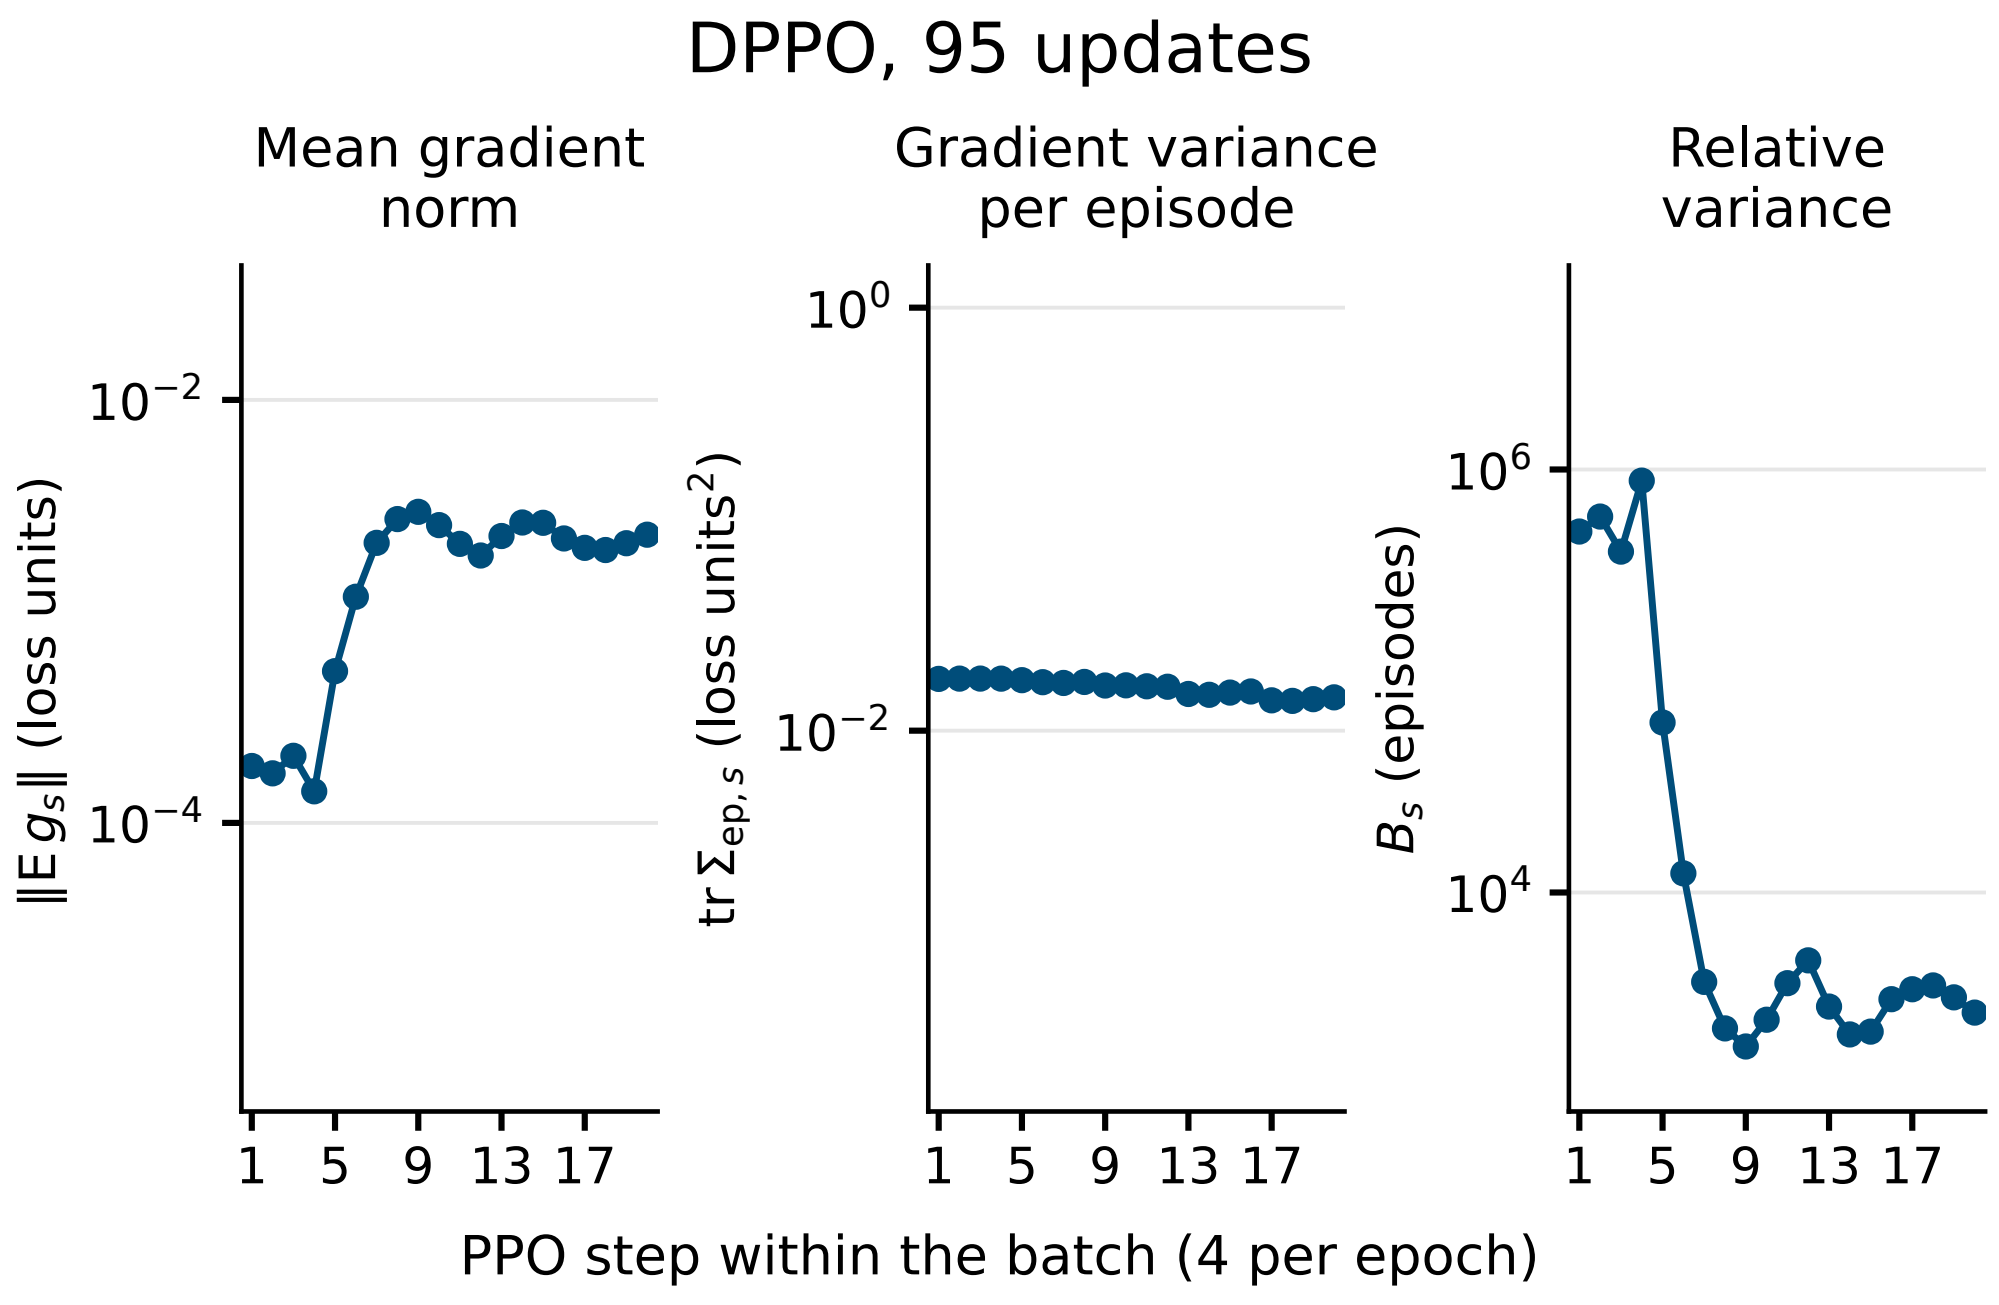
DPPO, 95 updates
Mean gradient norm
Gradient variance per episode
Relative variance
$\|\mathrm{E}\, g_s\|$ (loss units)
$\mathrm{tr}\, \Sigma_{\mathrm{ep},s}$ (loss units^2)
B_s (episodes)
10^{-2}
10^{-4}
10^{0}
10^{-2}
10^{6}
10^{4}
1 5 9 13 17
1 5 9 13 17
1 5 9 13 17
PPO step within the batch (4 per epoch)