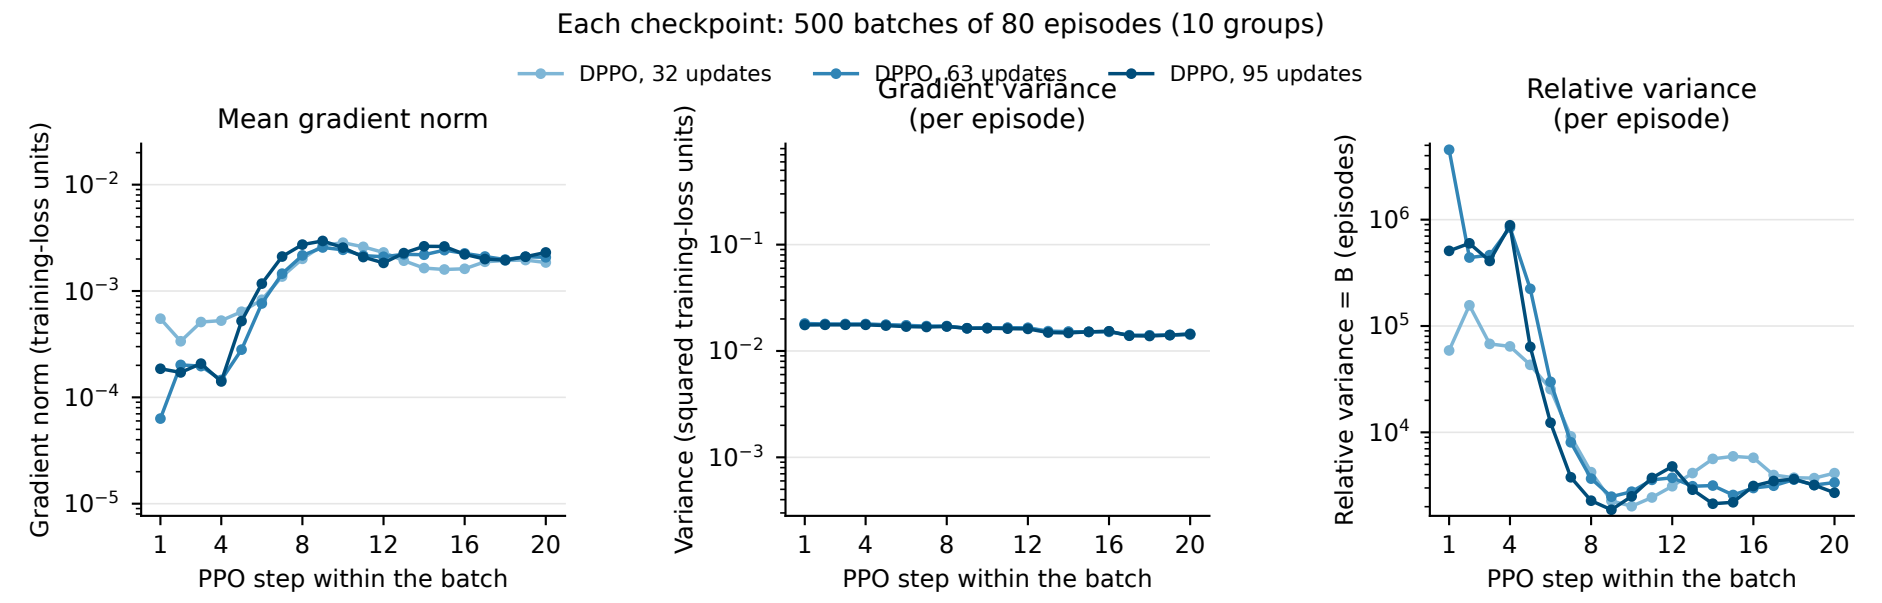
Each checkpoint: 500 batches of 80 episodes (10 groups)
DPPO, 32 updates
DPPO, 63 updates
DPPO, 95 updates
Mean gradient norm
Gradient norm (training-loss units)
PPO step within the batch
Gradient variance
(per episode)
Variance (squared training-loss units)
PPO step within the batch
Relative variance
(per episode)
Relative variance = B (episodes)
PPO step within the batch