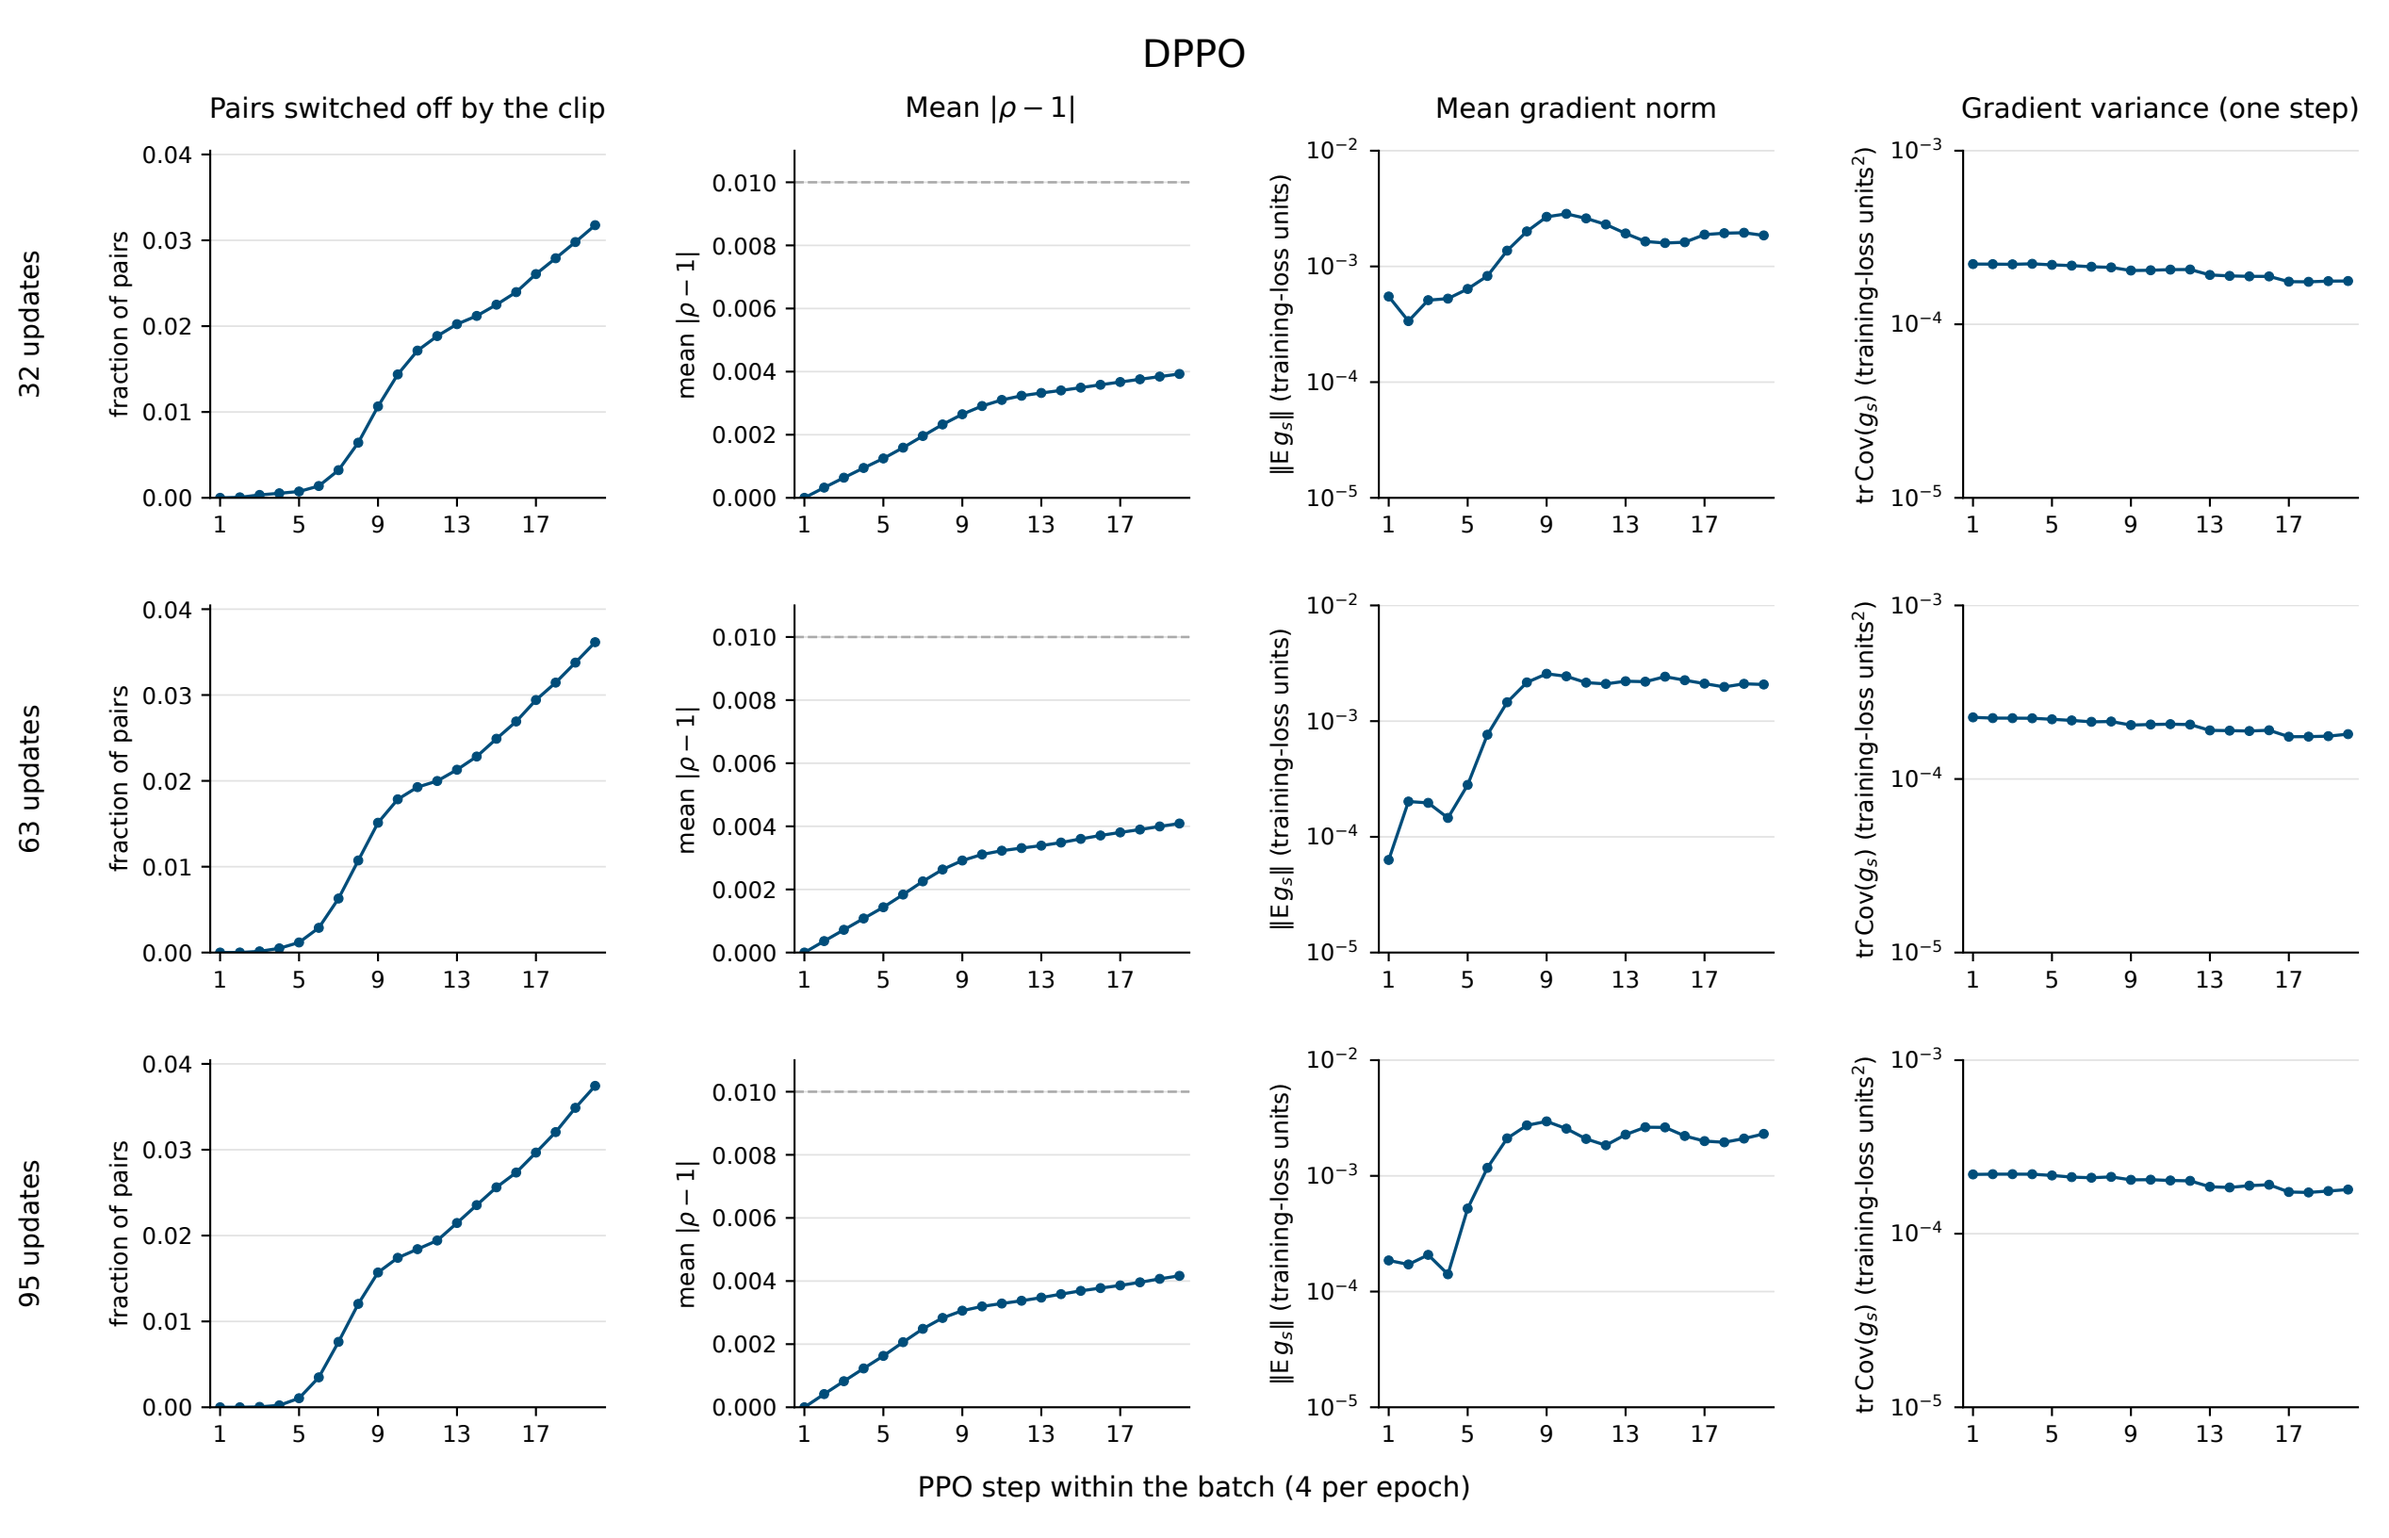
DPPO
Pairs switched off by the clip
Mean $|\rho - 1|$
Mean gradient norm
Gradient variance (one step)
32 updates
63 updates
95 updates
fraction of pairs
mean $|\rho - 1|$
$\|\mathrm{E}\, g_s\|$ (training-loss units)
$\mathrm{tr}\,\mathrm{Cov}(g_s)$ (training-loss units^2)
PPO step within the batch (4 per epoch)