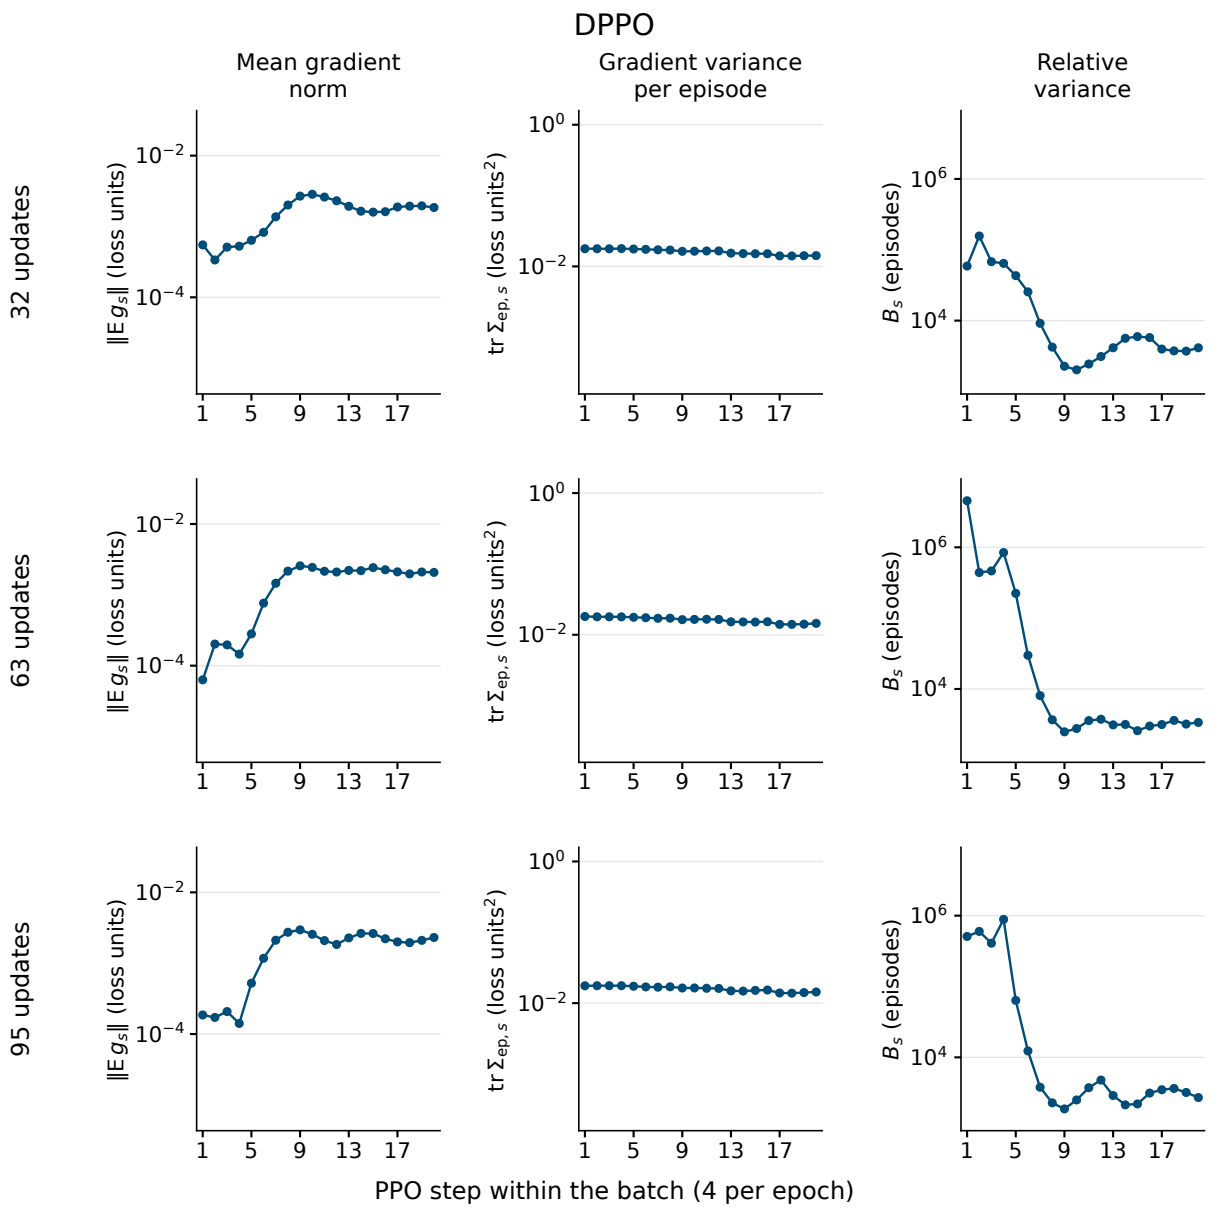
DPPO
Mean gradient norm
Gradient variance per episode
Relative variance
32 updates
63 updates
95 updates
$\|\mathbb{E}\, g_s\|$ (loss units)
$\operatorname{tr} \Sigma_{\text{ep},s}$ (loss units^2)
B_s (episodes)
10^{-2}
10^{-4}
10^{0}
10^{6}
10^{4}
1
5
9
13
17
PPO step within the batch (4 per epoch)