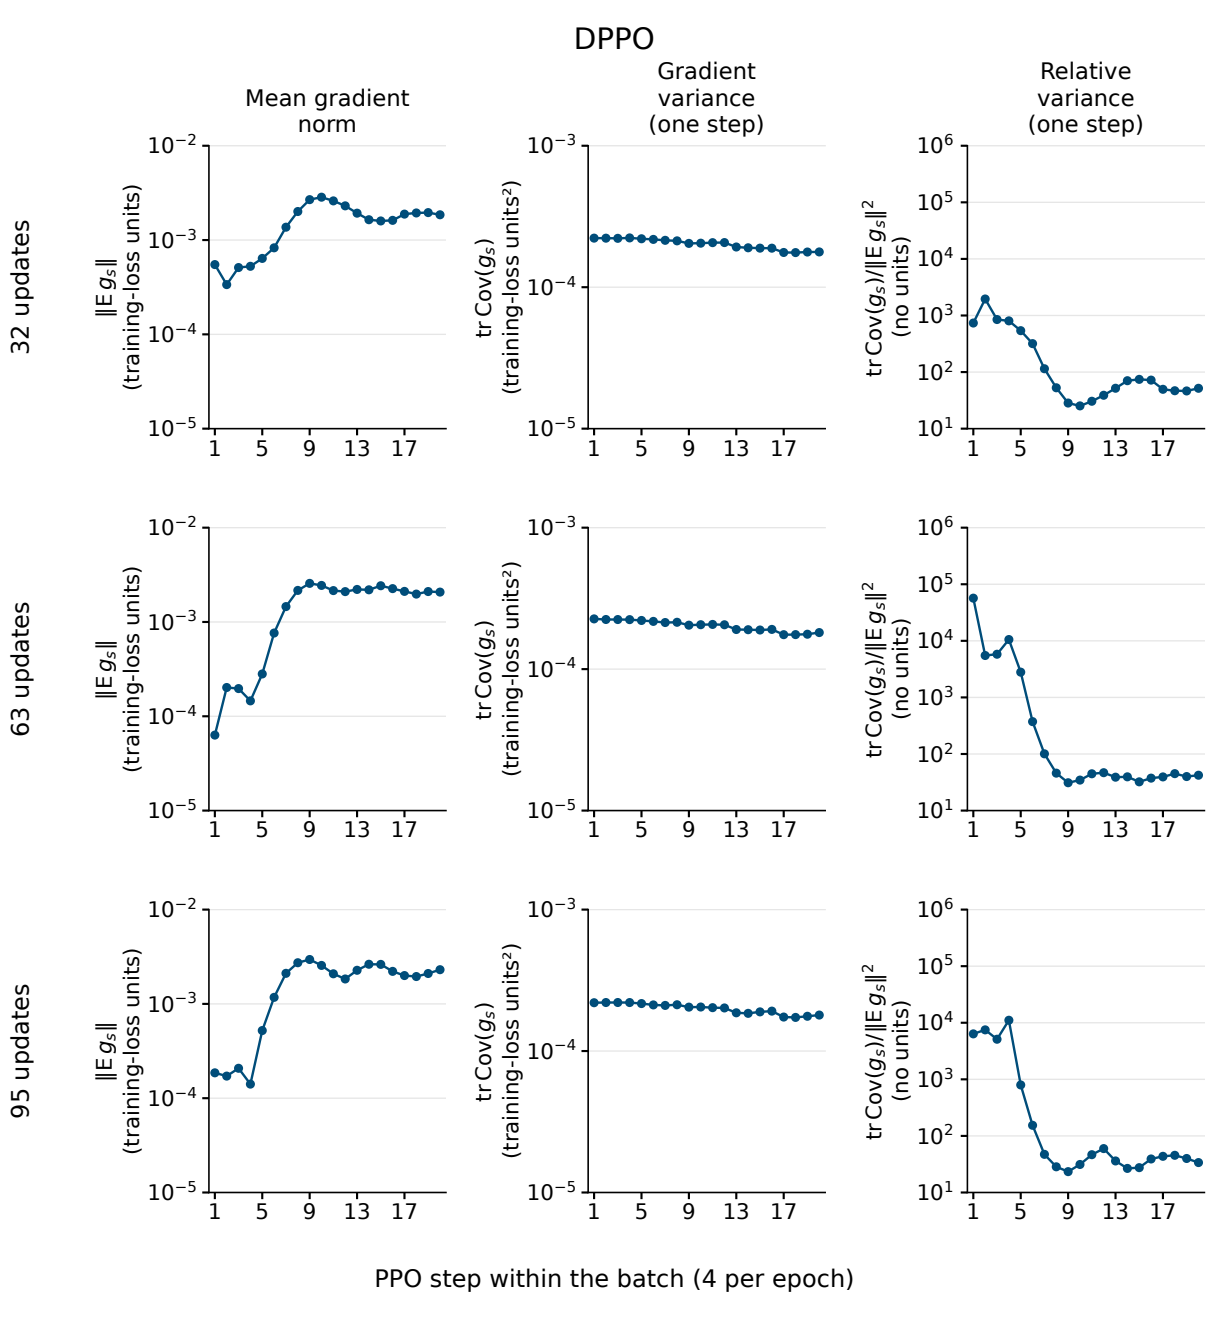
DPPO
Mean gradient norm
Gradient variance (one step)
Relative variance (one step)
32 updates
63 updates
95 updates
$\|\mathrm{E}\, g_s\|$ (training-loss units)
$\mathrm{tr}\,\mathrm{Cov}(g_s)$ (training-loss units$^2$)
$\mathrm{tr}\,\mathrm{Cov}(g_s)/\|\mathrm{E}\, g_s\|^2$ (no units)
PPO step within the batch (4 per epoch)