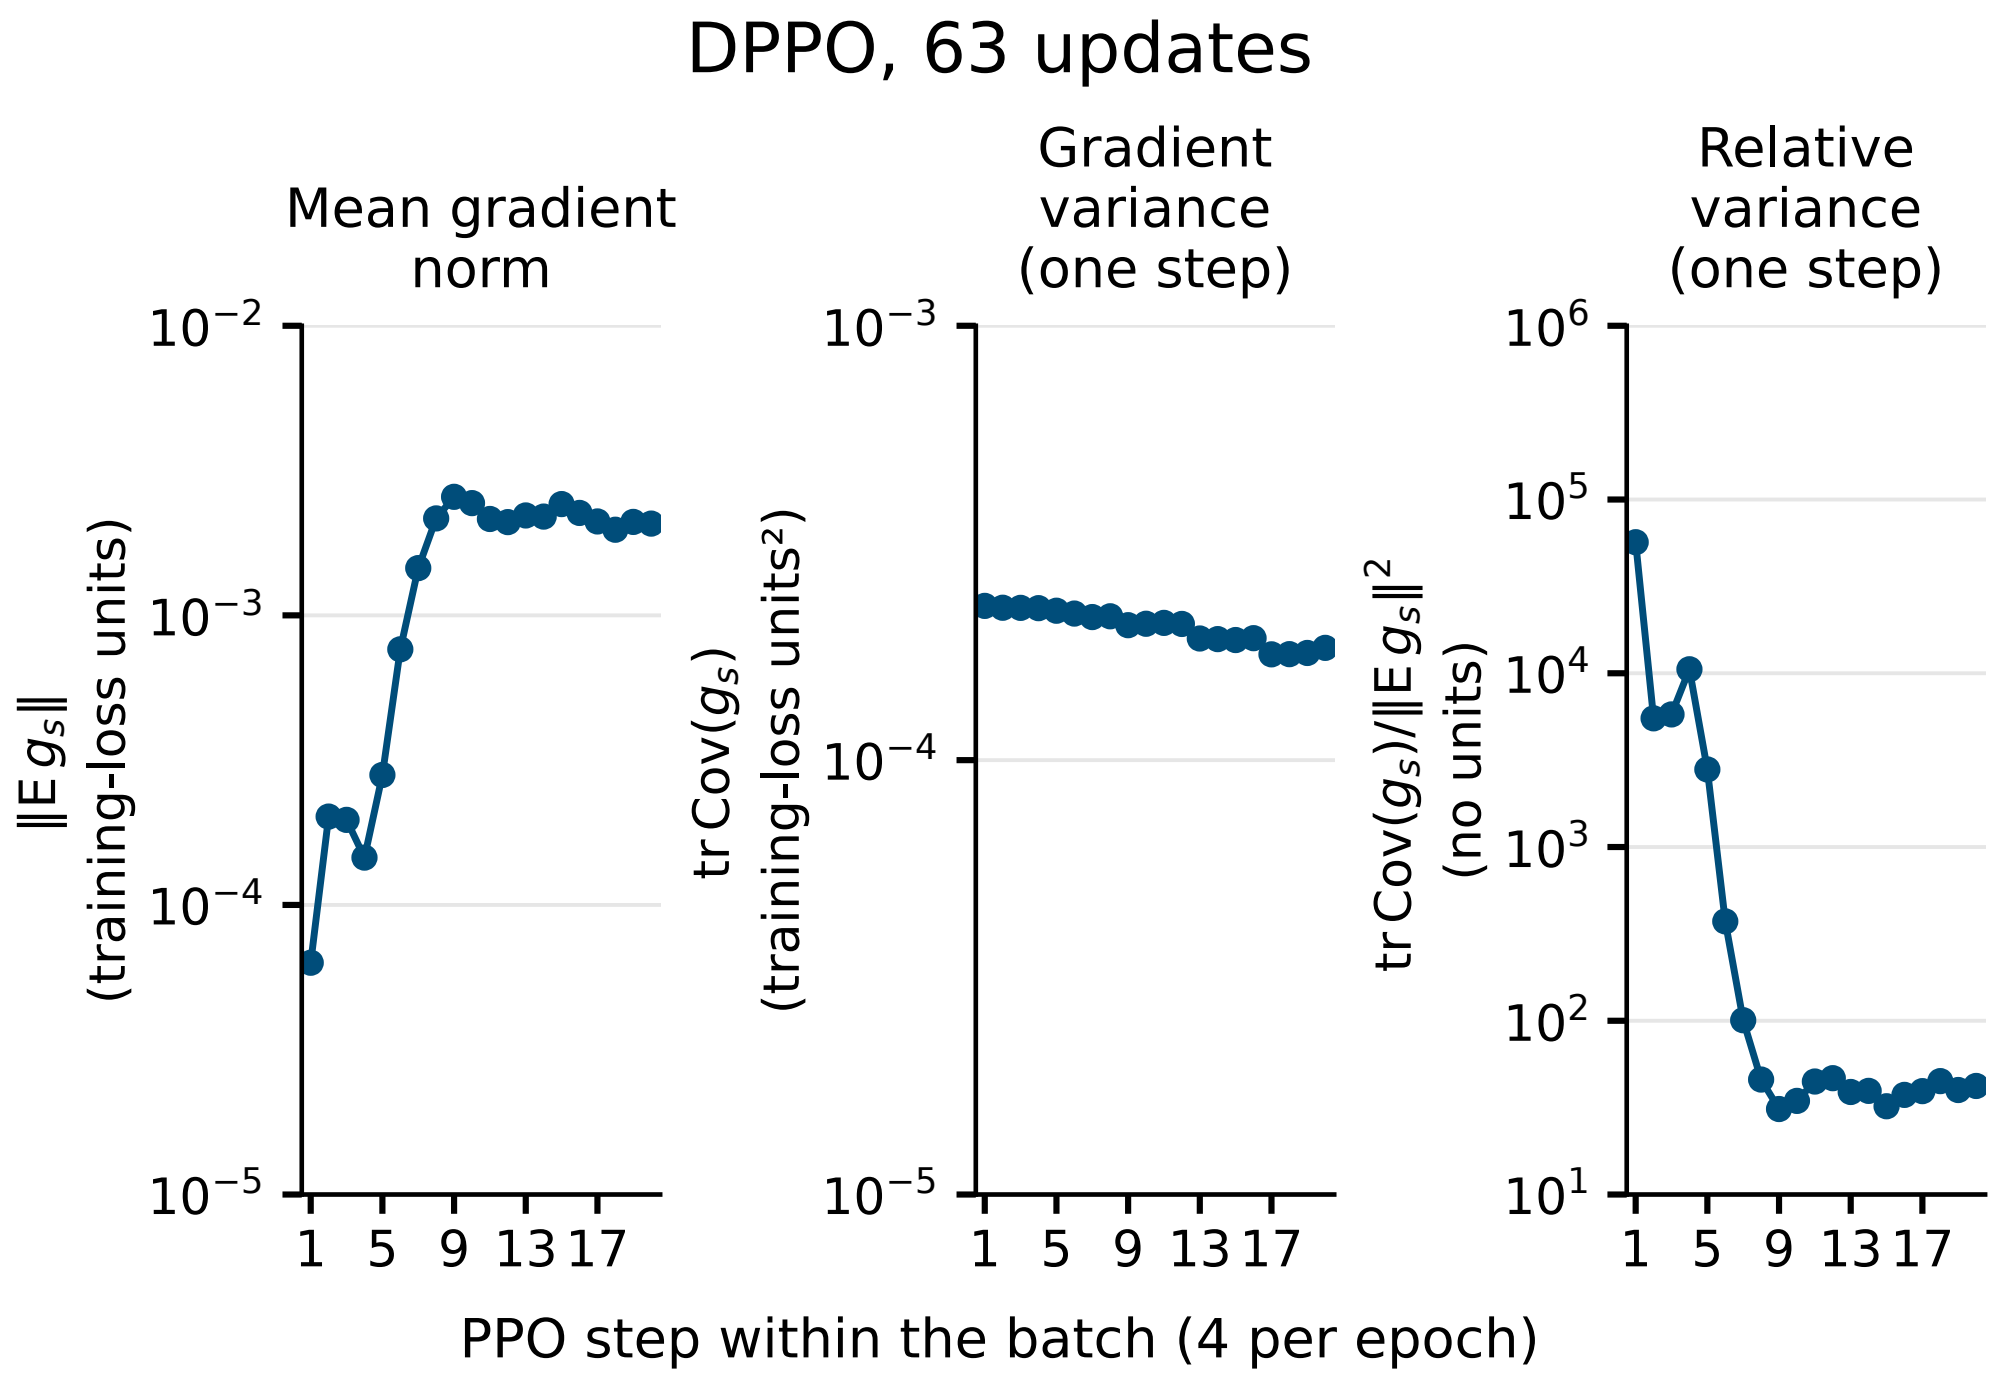
DPPO, 63 updates
Mean gradient norm
Gradient variance (one step)
Relative variance (one step)
$\|\mathrm{E}\, g_s\|$ (training-loss units)
$\mathrm{tr}\,\mathrm{Cov}(g_s)$ (training-loss units²)
$\mathrm{tr}\,\mathrm{Cov}(g_s)/\|\mathrm{E}\, g_s\|^2$ (no units)
$10^{-2}$
$10^{-3}$
$10^{-4}$
$10^{-5}$
$10^{6}$
$10^{5}$
$10^{4}$
$10^{3}$
$10^{2}$
$10^{1}$
1 5 9 13 17
PPO step within the batch (4 per epoch)