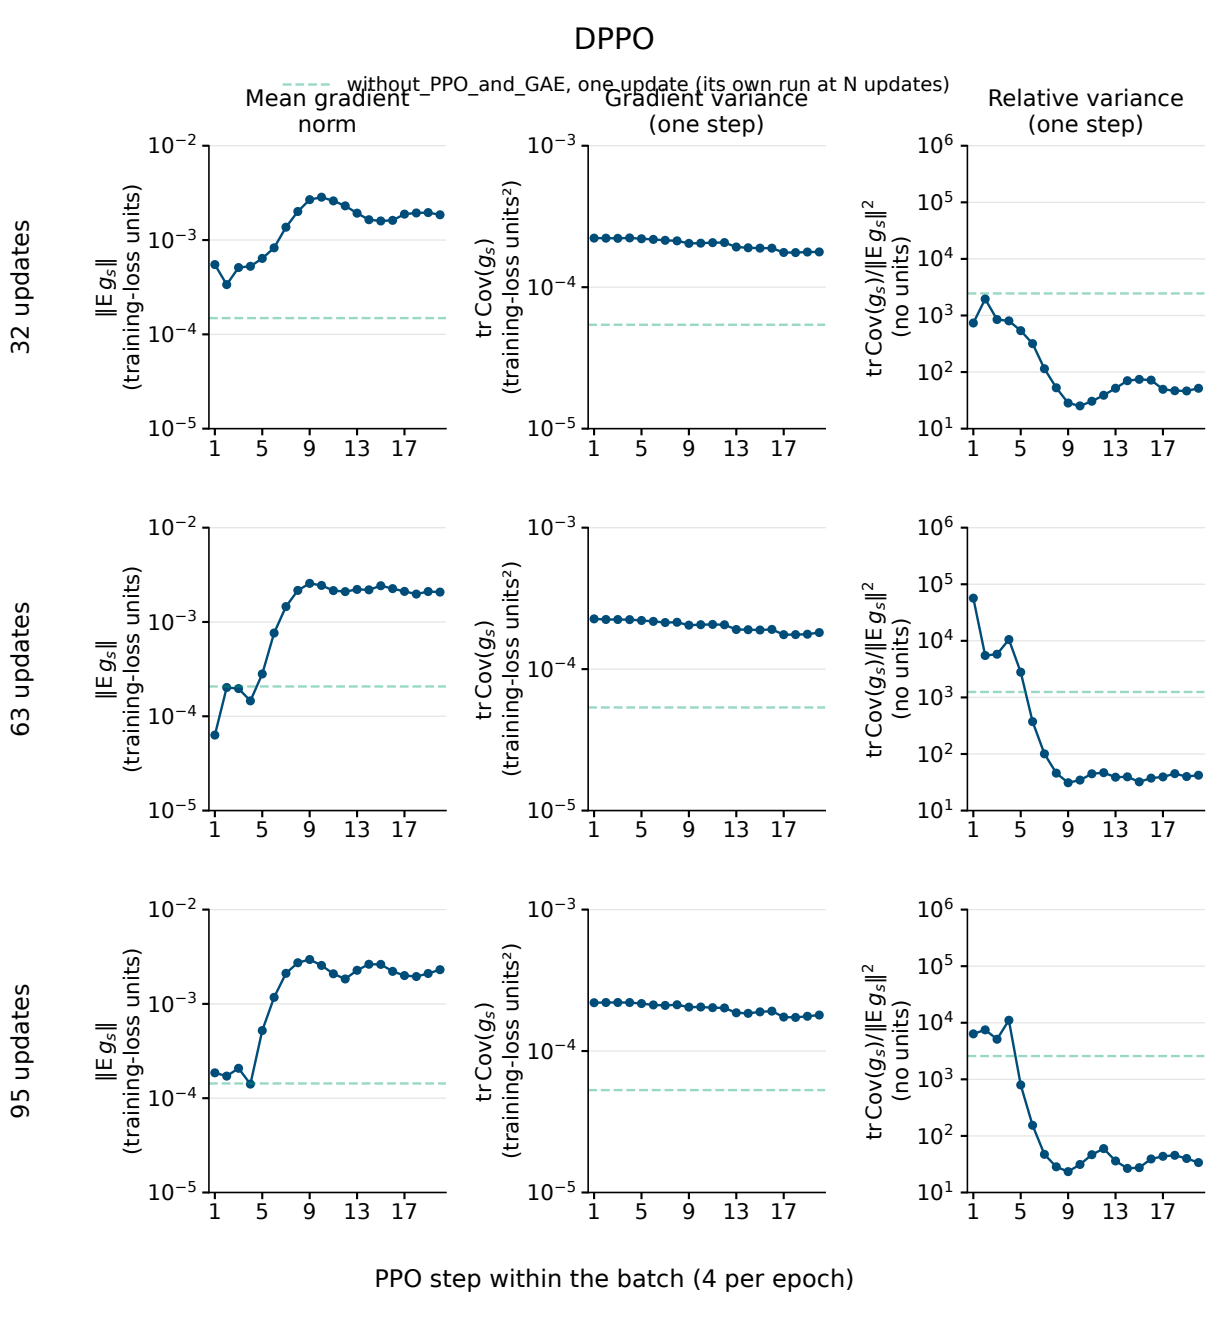
DPPO
without_PPO_and_GAE, one update (its own run at N updates)
Mean gradient norm
Gradient variance (one step)
Relative variance (one step)
32 updates
63 updates
95 updates
$\|\mathrm{E}\, g_s\|$ (training-loss units)
$\mathrm{tr}\,\mathrm{Cov}(g_s)$ (training-loss units²)
$\mathrm{tr}\,\mathrm{Cov}(g_s)/\|\mathrm{E}\, g_s\|^2$ (no units)
PPO step within the batch (4 per epoch)