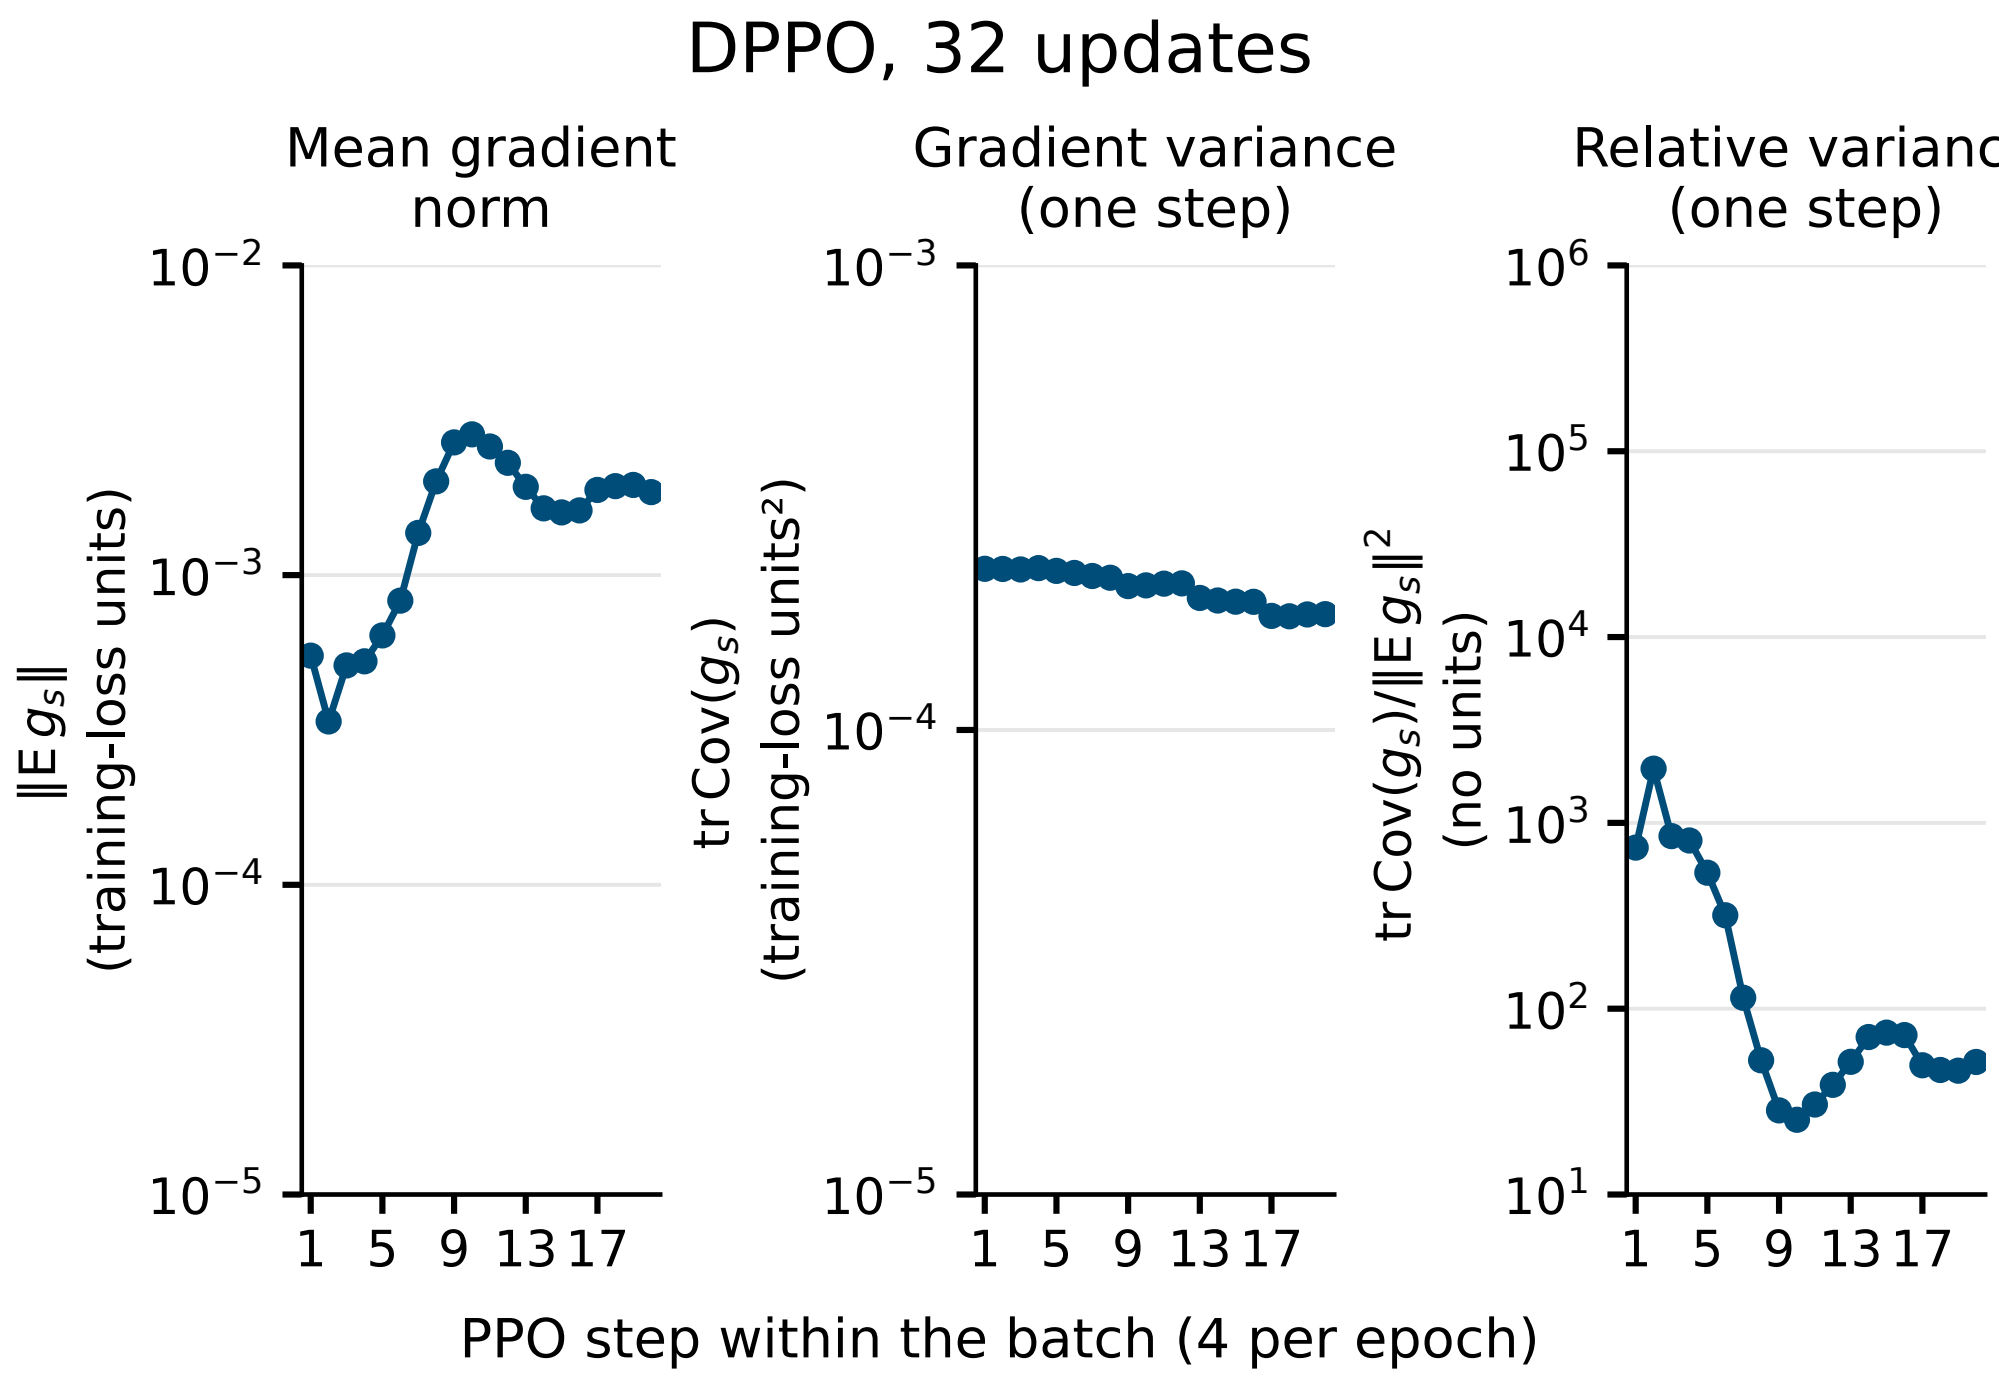
DPPO, 32 updates
Mean gradient norm
Gradient variance (one step)
Relative variance (one step)
$\|\mathrm{E}\, g_s\|$ (training-loss units)
$\mathrm{tr}\,\mathrm{Cov}(g_s)$ (training-loss units²)
$\mathrm{tr}\,\mathrm{Cov}(g_s)/\|\mathrm{E}\, g_s\|^2$ (no units)
10^{-2}
10^{-3}
10^{-4}
10^{-5}
10^{-3}
10^{-4}
10^{-5}
10^{6}
10^{5}
10^{4}
10^{3}
10^{2}
10^{1}
1 5 9 13 17
1 5 9 13 17
1 5 9 13 17
PPO step within the batch (4 per epoch)