


Phase 1: provisional gradient-noise point estimates


NO VALID CONFIDENCE INTERVALS — existing percentile intervals and positivity claims are unreliable

● Pretrained policy θ_0 ■ NC4 1e-4 final checkpoint

Critical batch size in episodes



Expected training-batch gradient cosine



N/A: nonpositive corrected squared-signal estimate; no zero or lower-bound substitution. No error bars are justified.

V1: scaled GAE (critic); V2: scaled MC – critic; V3: raw MC – LOO. V1-V3 normalize coefficients.

V4: raw MC – LOO with decision discount; V5: raw MC with decision discount. V4-V5 do not normalize.

Per policy: 25 fresh batches $\times$ 40 environments = 1,000 per-environment gradients; two episodes per environment rollout.